

DBA3702 Assignment 1

Data Wrangling & Social Network Analysis - TechConnect Inc.

Group X - [Names and Student IDs]

2026-01-24

Contents

1 Part 1: Data Wrangling with dplyr (45 points)	3
1.1 Question 1.1: Data Exploration (5 points)	3
1.1.1 a) Load packages and read data	3
1.1.2 b) Convert to tibble and display first 10 rows	3
1.1.3 c) Data summary	3
1.2 Question 1.2: Selecting and Filtering (8 points)	5
1.2.1 d) Select specific columns	5
1.2.2 e) Filter performance > 4.0	5
1.2.3 f) Engineering/Marketing with > 5 years experience	6
1.2.4 g) Select using helpers	6
1.3 Question 1.3: Sorting and Ranking (7 points)	8
1.3.1 h) Top 5 highest-paid	8
1.3.2 i) Sort by department then performance	8
1.3.3 j) Lowest salary in each department	8
1.4 Question 1.4: Creating New Variables (10 points)	10
1.4.1 k) Salary per year of experience	10
1.4.2 l) Performance category	10
1.4.3 m) Experience level	11
1.4.4 n) High performer flag	12
1.5 Question 1.5: Aggregation and Grouping (15 points)	13
1.5.1 o) Company-wide summary	13
1.5.2 p) Summary by department	13
1.5.3 q) Summary by department and role	13

1.5.4	r) Salary as % of department average	14
1.5.5	s) Top 3 departments by performance (3+ years exp)	15
2	Part 2: Social Network Analysis (45 points)	16
2.1	Question 2.1: Network Construction and Visualization (10 points)	16
2.1.1	t) Load network data	16
2.1.2	u) Construct undirected graph	16
2.1.3	v) Basic and improved plots	17
2.1.4	w) Department-colored network	17
2.2	Question 2.2: Connected Components (8 points)	19
2.2.1	x) Find connected components	19
2.2.2	y) Largest component size	19
2.2.3	z) Extract and plot largest component	19
2.2.4	aa) Why use the largest connected component?	20
2.3	Question 2.3: Centrality Metrics (15 points)	21
2.3.1	bb) Degree Centrality	21
2.3.2	cc) Closeness Centrality	22
2.3.3	dd) Betweenness Centrality	23
2.3.4	ee) PageRank	23
2.3.5	ff) Comparing all metrics	24
2.4	Question 2.4: Community Detection (12 points)	26
2.4.1	gg) Spinglass clustering	26
2.4.2	hh) Community sizes	26
2.4.3	ii) Visualize by community	26
2.4.4	jj) Community vs Department	27
2.4.5	kk) Business insights	28
3	Part 3: Integration and Insights (10 points)	29
3.1	Question 3.1: Joining Data (5 points)	29
3.1.1	ll) Join employee data with centrality metrics	29
3.1.2	mm) Correlation analysis	29
3.1.3	nn) High performers with low centrality	30
3.2	Question 3.2: Executive Summary (5 points)	31
3.2.1	Executive Summary for TechConnect Management	31

1 Part 1: Data Wrangling with dplyr (45 points)

1.1 Question 1.1: Data Exploration (5 points)

1.1.1 a) Load packages and read data

```
library(dplyr)
library(tibble)

employees <- read.csv("data/employees.csv")
```

1.1.2 b) Convert to tibble and display first 10 rows

```
employees <- as_tibble(employees)
print(employees, n = 10)
```

```
## # A tibble: 50 x 8
##   employee_id name      department role    years_exp salary performance_score
##       <int> <chr>     <chr>     <chr>     <int> <int>          <dbl>
## 1           1 Alice Chen  Engineering Senior      8 95000        4.5
## 2           2 Bob Martinez Engineering Lead       12 120000       4.8
## 3           3 Charlie Kim  Engineering Junior     2 65000        3.8
## 4           4 Diana Patel Engineering Senior     7 92000        4.2
## 5           5 Eve Thompson Engineering Manager  15 140000       4.6
## 6           6 Frank Liu   Engineering Junior     1 58000        3.2
## 7           7 Grace Okonkwo Engineering Senior   9 98000        4.4
## 8           8 Henry Wang   Marketing Lead       10 105000       4.3
## 9           9 Iris Nakamura Marketing Senior    6 82000        3.9
## 10          10 Jack Brown  Marketing Junior     2 55000        3.5
## # i 40 more rows
## # i 1 more variable: projects_completed <int>
```

1.1.3 c) Data summary

```
cat("Rows:", nrow(employees), "\n")
```

```
## Rows: 50
```

```
cat("Columns:", ncol(employees), "\n")
```

```
## Columns: 8
```

```
sapply(employees, class)
```

```
##      employee_id             name      department        role
##      "integer"       "character"     "character"     "character"
##      years_exp          salary  performance_score projects_completed
##      "integer"       "integer"       "numeric"       "integer"
```

```
glimpse(employees)
```

```
## Rows: 50
## Columns: 8
## $ employee_id      <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17~
## $ name            <chr> "Alice Chen", "Bob Martinez", "Charlie Kim", "Diana Patel~
## $ department      <chr> "Engineering", "Engineering", "Engineering", "Engineering~
## $ role            <chr> "Senior", "Lead", "Junior", "Senior", "Manager", "Junior"~
## $ years_exp       <int> 8, 12, 2, 7, 15, 1, 9, 10, 6, 2, 14, 5, 1, 11, 8, 3, 6, 1~
## $ salary          <int> 95000, 120000, 65000, 92000, 140000, 58000, 98000, 105000~
## $ performance_score <dbl> 4.5, 4.8, 3.8, 4.2, 4.6, 3.2, 4.4, 4.3, 3.9, 3.5, 4.5, 4.~
## $ projects_completed <int> 15, 22, 5, 12, 28, 3, 16, 18, 11, 4, 24, 9, 2, 21, 14, 6,~
```

```
summary(employees)
```

```
##   employee_id      name      department        role
##   Min. : 1.00  Length:50      Length:50      Length:50
##   1st Qu.:13.25 Class :character  Class :character  Class :character
##   Median :25.50 Mode  :character  Mode  :character  Mode  :character
##   Mean   :25.50
##   3rd Qu.:37.75
##   Max.   :50.00
## 
##   years_exp          salary  performance_score projects_completed
##   Min.   : 1.00  Min.   :48000  Min.   :2.900  Min.   : 2.00
##   1st Qu.: 3.00  1st Qu.:65750  1st Qu.:3.600  1st Qu.: 6.00
##   Median : 6.50  Median :85000  Median :4.000  Median :11.50
##   Mean   : 6.68  Mean   :86420  Mean   :3.970  Mean   :12.24
##   3rd Qu.: 9.00  3rd Qu.:101500 3rd Qu.:4.375  3rd Qu.:17.00
##   Max.   :16.00  Max.   :145000  Max.   :4.900  Max.   :30.00
```

We have 50 employees with 8 variables. The data includes basic info like name and department, plus numeric stuff like salary, years of experience, performance scores, and project counts. Performance ranges from around 2.9 to 4.9, and experience goes from 1 to 16 years.

1.2 Question 1.2: Selecting and Filtering (8 points)

1.2.1 d) Select specific columns

```
employees %>%
  select(name, department, role, performance_score)
```

```
## # A tibble: 50 x 4
##   name      department  role  performance_score
##   <chr>     <chr>       <chr>          <dbl>
## 1 Alice Chen  Engineering Senior        4.5
## 2 Bob Martinez Engineering Lead        4.8
## 3 Charlie Kim Engineering Junior      3.8
## 4 Diana Patel Engineering Senior      4.2
## 5 Eve Thompson Engineering Manager    4.6
## 6 Frank Liu   Engineering Junior      3.2
## 7 Grace Okonkwo Engineering Senior    4.4
## 8 Henry Wang   Marketing Lead        4.3
## 9 Iris Nakamura Marketing Senior      3.9
## 10 Jack Brown  Marketing Junior      3.5
## # i 40 more rows
```

1.2.2 e) Filter performance > 4.0

```
high_performers <- employees %>%
  filter(performance_score > 4.0)
high_performers
```

```
## # A tibble: 23 x 8
##   employee_id name      department  role  years_exp salary performance_score
##   <int> <chr>     <chr>       <chr> <int> <int>          <dbl>
## 1 1 Alice Chen  Engineering Senior 8 95000 4.5
## 2 2 Bob Martinez Engineering Lead 12 120000 4.8
## 3 4 Diana Patel Engineering Senior 7 92000 4.2
## 4 5 Eve Thompson Engineering Manager 15 140000 4.6
## 5 7 Grace Okonkwo Engineering Senior 9 98000 4.4
## 6 8 Henry Wang   Marketing Lead 10 105000 4.3
## 7 11 Kate Wilson Marketing Manager 14 125000 4.5
## 8 14 Nathan Lee  Sales Lead 11 115000 4.7
## 9 15 Olivia Davis Sales Senior 8 88000 4.1
## 10 18 Rachel Green Sales Manager 13 130000 4.4
## # i 13 more rows
## # i 1 more variable: projects_completed <int>
```

```
cat("Count:", nrow(high_performers))
```

Count: 23

1.2.3 f) Engineering/Marketing with > 5 years experience

```
employees %>%
  filter((department == "Engineering" | department == "Marketing") & years_exp > 5)

## # A tibble: 14 x 8
##   employee_id name      department role    years_exp salary performance_score
##   <int> <chr>     <chr>    <chr>    <int> <int>          <dbl>
## 1 1       Alice Chen  Engineering Senior     8 95000        4.5
## 2 2       Bob Martinez Engineering Lead      12 120000       4.8
## 3 4       Diana Patel Engineering Senior     7 92000        4.2
## 4 5       Eve Thompson Engineering Manager  15 140000       4.6
## 5 7       Grace Okonkwo Engineering Senior   9 98000        4.4
## 6 8       Henry Wang   Marketing Lead      10 105000       4.3
## 7 9       Iris Nakamura Marketing Senior    6 82000        3.9
## 8 11      Kate Wilson  Marketing Manager   14 125000       4.5
## 9 31      Eric Zhang   Engineering Senior   6 88000        4
## 10 33     George Park  Marketing Senior    7 80000        3.7
## 11 41     Oscar Rivera  Engineering Lead    11 118000       4.5
## 12 42     Paula Hughes  Engineering Manager 16 145000       4.9
## 13 43     Quentin Price Marketing Lead     8 100000       4.1
## 14 47     Ulrich Weber  Engineering Senior  7 94000        4.2
## # i 1 more variable: projects_completed <int>
```

1.2.4 g) Select using helpers

```
employees %>%
  select(contains("score") | starts_with("p"))
```

```
## # A tibble: 50 x 2
##   performance_score projects_completed
##   <dbl>           <int>
## 1 4.5              15
## 2 4.8              22
## 3 3.8              5
## 4 4.2              12
## 5 4.6              28
## 6 3.2              3
```

```
## 7          4.4          16
## 8          4.3          18
## 9          3.9          11
## 10         3.5           4
## # i 40 more rows
```

This grabs performance_score and projects_completed - the columns with “score” in the name or starting with “p”.

1.3 Question 1.3: Sorting and Ranking (7 points)

1.3.1 h) Top 5 highest-paid

```
employees %>%
  arrange(desc(salary)) %>%
  head(5)

## # A tibble: 5 x 8
##   employee_id name      department role    years_exp salary performance_score
##       <int> <chr>     <chr>     <chr>     <int>   <int>           <dbl>
## 1        42 Paula Hughes Engineering Manager      16 145000        4.9
## 2         5 Eve Thompson Engineering Manager      15 140000        4.6
## 3        28 Bella Moore  Finance    Manager      14 135000        4.7
## 4        18 Rachel Green Sales      Manager      13 130000        4.4
## 5        11 Kate Wilson Marketing Manager      14 125000        4.5
## # i 1 more variable: projects_completed <int>
```

1.3.2 i) Sort by department then performance

```
employees %>%
  arrange(department, desc(performance_score))

## # A tibble: 50 x 8
##   employee_id name      department role    years_exp salary performance_score
##       <int> <chr>     <chr>     <chr>     <int>   <int>           <dbl>
## 1        42 Paula Hughes Engineering Manager      16 145000        4.9
## 2         2 Bob Martinez Engineering Lead        12 120000        4.8
## 3         5 Eve Thompson Engineering Manager      15 140000        4.6
## 4         1 Alice Chen   Engineering Senior       8  95000        4.5
## 5        41 Oscar Rivera Engineering Lead       11 118000        4.5
## 6         7 Grace Okonkwo Engineering Senior      9  98000        4.4
## 7         4 Diana Patel  Engineering Senior       7  92000        4.2
## 8        47 Ulrich Weber Engineering Senior       7  94000        4.2
## 9         31 Eric Zhang   Engineering Senior      6  88000        4
## 10        32 Fiona O'Brien Engineering Junior     3  68000        3.9
## # i 40 more rows
## # i 1 more variable: projects_completed <int>
```

1.3.3 j) Lowest salary in each department

```
employees %>%
  arrange(department, salary) %>%
  group_by(department) %>%
  slice_head(n = 1) %>%
  ungroup()

## # A tibble: 5 x 8
##   employee_id name      department role    years_exp salary performance_score
##       <int> <chr>     <chr>     <chr>     <int> <int>          <dbl>
## 1           6 Frank Liu  Engineering Junior      1  58000        3.2
## 2          30 Dana Hill Finance    Junior      1  55000        3
## 3          22 Victor Nguyen HR        Junior      1  48000        3.1
## 4          13 Maya Rodriguez Marketing Junior      1  52000        3.3
## 5          36 Julia Foster  Sales      Junior      1  53000        2.9
## # i 1 more variable: projects_completed <int>
```

1.4 Question 1.4: Creating New Variables (10 points)

1.4.1 k) Salary per year of experience

```
employees %>%
  mutate(salary_per_year_exp = salary / years_exp) %>%
  select(name, salary, years_exp, salary_per_year_exp)

## # A tibble: 50 x 4
##   name      salary years_exp salary_per_year_exp
##   <chr>     <int>    <int>          <dbl>
## 1 Alice Chen  95000       8        11875
## 2 Bob Martinez 120000      12       10000
## 3 Charlie Kim  65000       2        32500
## 4 Diana Patel  92000       7       13143.
## 5 Eve Thompson 140000      15       9333.
## 6 Frank Liu    58000       1        58000
## 7 Grace Okonkwo 98000       9       10889.
## 8 Henry Wang   105000      10       10500
## 9 Iris Nakamura 82000       6       13667.
## 10 Jack Brown  55000       2        27500
## # i 40 more rows
```

1.4.2 l) Performance category

```
employees_cat <- employees %>%
  mutate(performance_category = case_when(
    performance_score >= 4.5 ~ "Outstanding",
    performance_score >= 3.5 ~ "Exceeds Expectations",
    performance_score >= 2.5 ~ "Meets Expectations",
    TRUE ~ "Needs Improvement"
  ))

employees_cat %>%
  select(name, performance_score, performance_category)
```

```
## # A tibble: 50 x 3
##   name      performance_score performance_category
##   <chr>           <dbl> <chr>
## 1 Alice Chen        4.5 Outstanding
## 2 Bob Martinez      4.8 Outstanding
## 3 Charlie Kim       3.8 Exceeds Expectations
## 4 Diana Patel       4.2 Exceeds Expectations
## 5 Eve Thompson      4.6 Outstanding
```

```

## 6 Frank Liu           3.2 Meets Expectations
## 7 Grace Okonkwo      4.4 Exceeds Expectations
## 8 Henry Wang          4.3 Exceeds Expectations
## 9 Iris Nakamura       3.9 Exceeds Expectations
## 10 Jack Brown         3.5 Exceeds Expectations
## # i 40 more rows

```

```
employees_cat %>% count(performance_category)
```

```

## # A tibble: 3 x 2
##   performance_category     n
##   <chr>                  <int>
## 1 Exceeds Expectations    31
## 2 Meets Expectations      9
## 3 Outstanding             10

```

1.4.3 m) Experience level

```

employees_exp <- employees %>%
  mutate(experience_level = case_when(
    years_exp <= 3 ~ "Entry",
    years_exp <= 7 ~ "Mid",
    years_exp <= 12 ~ "Senior",
    TRUE ~ "Expert"
  ))
employees_exp %>%
  select(name, years_exp, experience_level)

```

```

## # A tibble: 50 x 3
##   name        years_exp experience_level
##   <chr>        <int> <chr>
## 1 Alice Chen      8 Senior
## 2 Bob Martinez   12 Senior
## 3 Charlie Kim     2 Entry
## 4 Diana Patel    7 Mid
## 5 Eve Thompson   15 Expert
## 6 Frank Liu       1 Entry
## 7 Grace Okonkwo   9 Senior
## 8 Henry Wang      10 Senior
## 9 Iris Nakamura    6 Mid
## 10 Jack Brown     2 Entry
## # i 40 more rows

```

```
employees_exp %>% count(experience_level)
```

```
## # A tibble: 4 x 2
##   experience_level     n
##   <chr>             <int>
## 1 Entry                 14
## 2 Expert                  5
## 3 Mid                   15
## 4 Senior                 16
```

1.4.4 n) High performer flag

```
employees %>%
  mutate(is_high_performer = performance_score > 4.0 & projects_completed >= 10) %>%
  filter(is_high_performer) %>%
  select(name, department, performance_score, projects_completed)
```

```
## # A tibble: 23 x 4
##   name      department  performance_score projects_completed
##   <chr>    <chr>            <dbl>                <int>
## 1 Alice Chen  Engineering       4.5                  15
## 2 Bob Martinez Engineering       4.8                  22
## 3 Diana Patel Engineering       4.2                  12
## 4 Eve Thompson Engineering       4.6                  28
## 5 Grace Okonkwo Engineering     4.4                  16
## 6 Henry Wang   Marketing        4.3                  18
## 7 Kate Wilson  Marketing        4.5                  24
## 8 Nathan Lee   Sales           4.7                  21
## 9 Olivia Davis  Sales           4.1                  14
## 10 Rachel Green Sales           4.4                  25
## # i 13 more rows
```

1.5 Question 1.5: Aggregation and Grouping (15 points)

1.5.1 o) Company-wide summary

```
employees %>%
  summarise(
    total_employees = n(),
    avg_salary = mean(salary),
    avg_performance = mean(performance_score),
    total_projects = sum(projects_completed)
  )

## # A tibble: 1 x 4
##   total_employees avg_salary avg_performance total_projects
##             <int>      <dbl>            <dbl>        <int>
## 1                 50       86420            3.97         612
```

1.5.2 p) Summary by department

```
employees %>%
  group_by(department) %>%
  summarise(
    count = n(),
    avg_salary = mean(salary),
    avg_perf = mean(performance_score),
    min_exp = min(years_exp),
    max_exp = max(years_exp)
  )

## # A tibble: 5 x 6
##   department  count avg_salary avg_perf min_exp max_exp
##   <chr>      <int>      <dbl>     <dbl>    <int>    <int>
## 1 Engineering    12      98417.    4.25      1       16
## 2 Finance        9      86333.    3.96      1       14
## 3 HR              9      75556.    3.76      1       12
## 4 Marketing      10      80700.    3.86      1       14
## 5 Sales           10      87600.    3.95      1       13
```

1.5.3 q) Summary by department and role

```
dept_role <- employees %>%
  group_by(department, role) %>%
  summarise(avg_salary = mean(salary), count = n(), .groups = "drop") %>%
```

```

arrange(desc(avg_salary))

dept_role

## # A tibble: 20 x 4
##   department role     avg_salary count
##   <chr>      <chr>     <dbl> <int>
## 1 Engineering Manager    142500     2
## 2 Finance     Manager    135000     1
## 3 Sales       Manager    130000     1
## 4 Marketing    Manager    125000     1
## 5 Engineering  Lead      119000     2
## 6 Sales       Lead      113500     2
## 7 HR          Manager    110000     1
## 8 Finance     Lead      105000     2
## 9 Marketing    Lead      102500     2
## 10 HR         Lead      93500      2
## 11 Engineering Senior    93400      5
## 12 Sales       Senior    87250      4
## 13 Finance     Senior    85667.     3
## 14 Marketing   Senior    79000      4
## 15 HR          Senior    71250      4
## 16 Engineering Junior   63667.     3
## 17 Finance     Junior   58333.     3
## 18 Sales       Junior   56667.     3
## 19 Marketing   Junior   53667.     3
## 20 HR          Junior   49000      2

```

```

# Highest combo:
dept_role %>% head(1)

```

```

## # A tibble: 1 x 4
##   department role     avg_salary count
##   <chr>      <chr>     <dbl> <int>
## 1 Engineering Manager    142500     2

```

The highest average salary is in the department-role combo shown above.

1.5.4 r) Salary as % of department average

```

employees %>%
  group_by(department) %>%
  mutate(
    dept_avg = mean(salary),

```

```

    pct_of_avg = salary / dept_avg * 100
) %>%
ungroup() %>%
arrange(desc(pct_of_avg)) %>%
select(name, department, salary, dept_avg, pct_of_avg)

```

```

## # A tibble: 50 x 5
##   name      department  salary dept_avg pct_of_avg
##   <chr>     <chr>     <int>    <dbl>      <dbl>
## 1 Bella Moore  Finance  135000  86333.     156.
## 2 Kate Wilson  Marketing 125000  80700      155.
## 3 Rachel Green Sales    130000  87600      148.
## 4 Paula Hughes Engineering 145000  98417.     147.
## 5 Wendy Clark   HR      110000  75556.     146.
## 6 Eve Thompson Engineering 140000  98417.     142.
## 7 Nathan Lee    Sales    115000  87600      131.
## 8 Henry Wang    Marketing 105000  80700      130.
## 9 Rosa Martinez Sales    112000  87600      128.
## 10 Tina White   HR      95000   75556.     126.
## # i 40 more rows

```

The person at the top earns the most relative to their department's average.

1.5.5 s) Top 3 departments by performance (3+ years exp)

```

employees %>%
filter(years_exp >= 3) %>%
group_by(department) %>%
summarise(avg_perf = mean(performance_score)) %>%
arrange(desc(avg_perf)) %>%
head(3)

```

```

## # A tibble: 3 x 2
##   department  avg_perf
##   <chr>        <dbl>
## 1 Engineering  4.4
## 2 Finance     4.3
## 3 Sales        4.15

```

2 Part 2: Social Network Analysis (45 points)

2.1 Question 2.1: Network Construction and Visualization (10 points)

2.1.1 t) Load network data

```
library(igraph)
library(RColorBrewer)

email_nodes <- read.csv("data/email_nodes.csv")
email_edges <- read.csv("data/email_edges.csv")

head(email_nodes)

##   id  department      role
## 1  1 Engineering Senior
## 2  2 Engineering    Lead
## 3  3 Engineering  Junior
## 4  4 Engineering Senior
## 5  5 Engineering Manager
## 6  6 Engineering  Junior

head(email_edges)

##   from to weight
## 1    1  2     25
## 2    1  3     15
## 3    1  4     20
## 4    1  5     30
## 5    1  7     18
## 6    2  3     22
```

2.1.2 u) Construct undirected graph

```
email_graph <- graph.data.frame(email_edges, vertices = email_nodes, directed = FALSE)

cat("Nodes:", vcount(email_graph), "\n")

## Nodes: 50

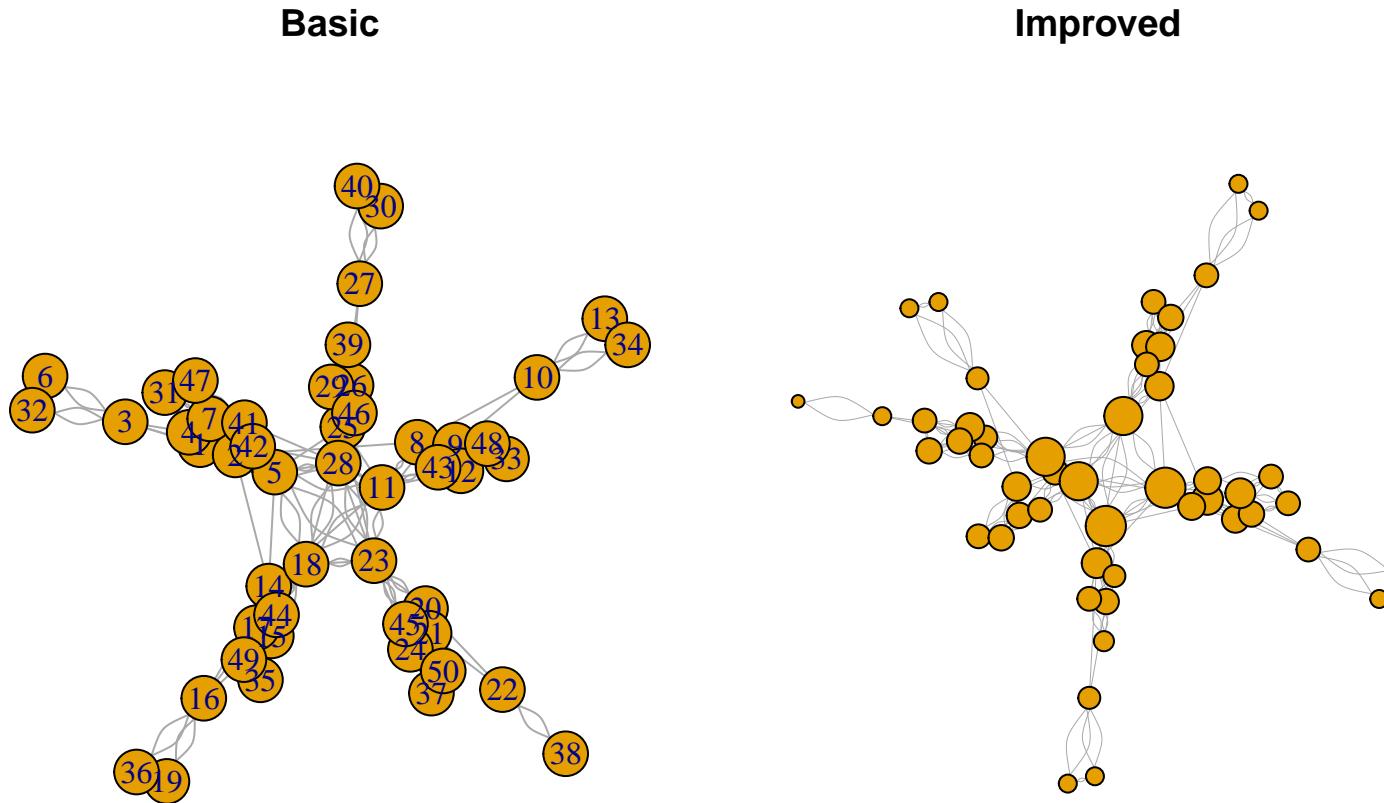
cat("Edges:", ecount(email_graph), "\n")

## Edges: 207
```

2.1.3 v) Basic and improved plots

```
par(mfrow = c(1, 2), mar = c(1, 1, 2, 1))
plot(email_graph, main = "Basic")

deg <- degree(email_graph)
plot(email_graph, vertex.label = NA, vertex.size = sqrt(deg) * 3,
     edge.width = 0.5, main = "Improved")
```



The improved version removes the messy labels and sizes nodes by how connected they are.

2.1.4 w) Department-colored network

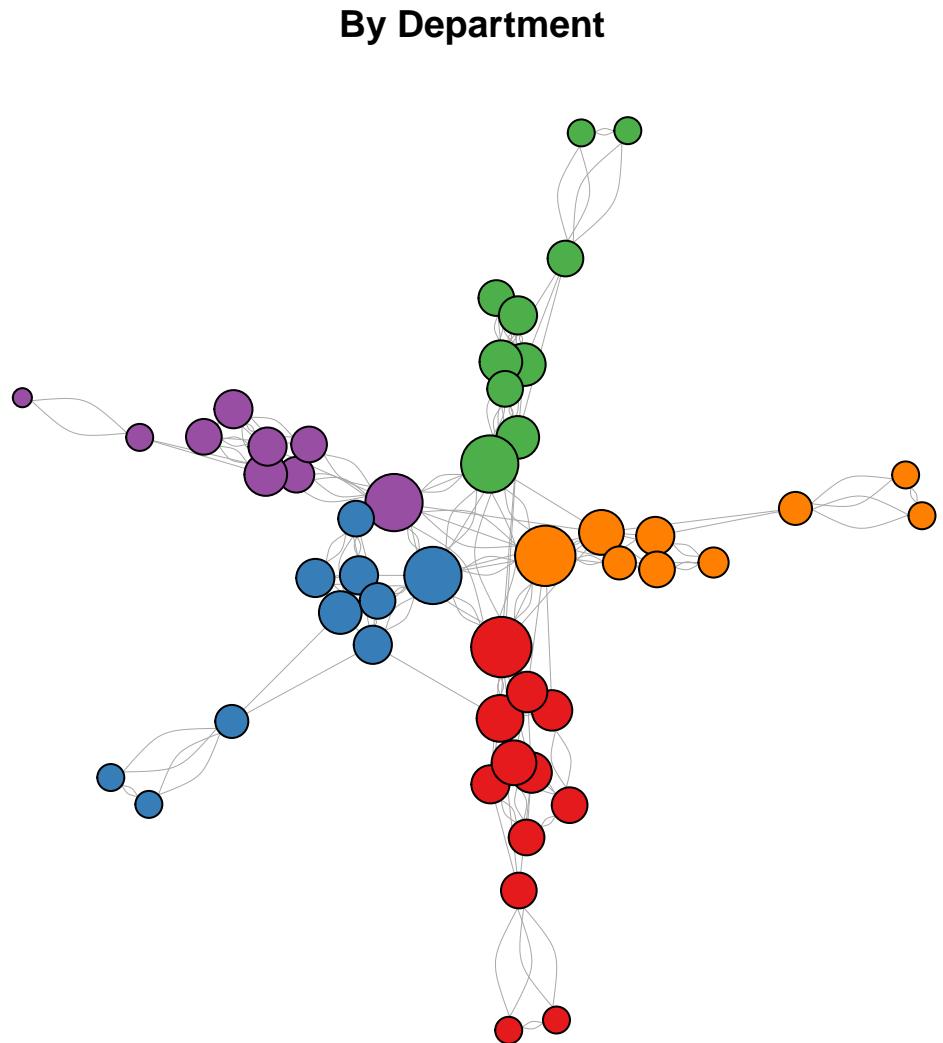
```
deps <- V(email_graph)$department
unique_depts <- unique(deps)
colors <- brewer.pal(length(unique_depts), "Set1")
```

```

names(colors) <- unique_depts

par(mfrow = c(1, 1), mar = c(1, 1, 2, 5))
plot(email_graph, vertex.label = NA, vertex.size = sqrt(deg) * 3,
     vertex.color = colors[depts], edge.width = 0.5,
     main = "By Department")
legend("topright", unique_depts, fill = colors, cex = 0.7, bty = "n")

```



2.2 Question 2.2: Connected Components (8 points)

2.2.1 x) Find connected components

```
comp <- components(email_graph)
cat("Number of components:", comp$no, "\n")
```

```
## Number of components: 1
```

2.2.2 y) Largest component size

```
lcc_size <- max(comp$csize)
cat("Largest component:", lcc_size, "employees\n")
```

```
## Largest component: 50 employees
```

```
cat("That's", round(lcc_size / vcount(email_graph) * 100, 1), "% of everyone\n")
```

```
## That's 100 % of everyone
```

2.2.3 z) Extract and plot largest component

```
lcc_id <- which.max(comp$csize)
lcc_nodes <- which(comp$membership == lcc_id)
lcc <- induced_subgraph(email_graph, lcc_nodes)

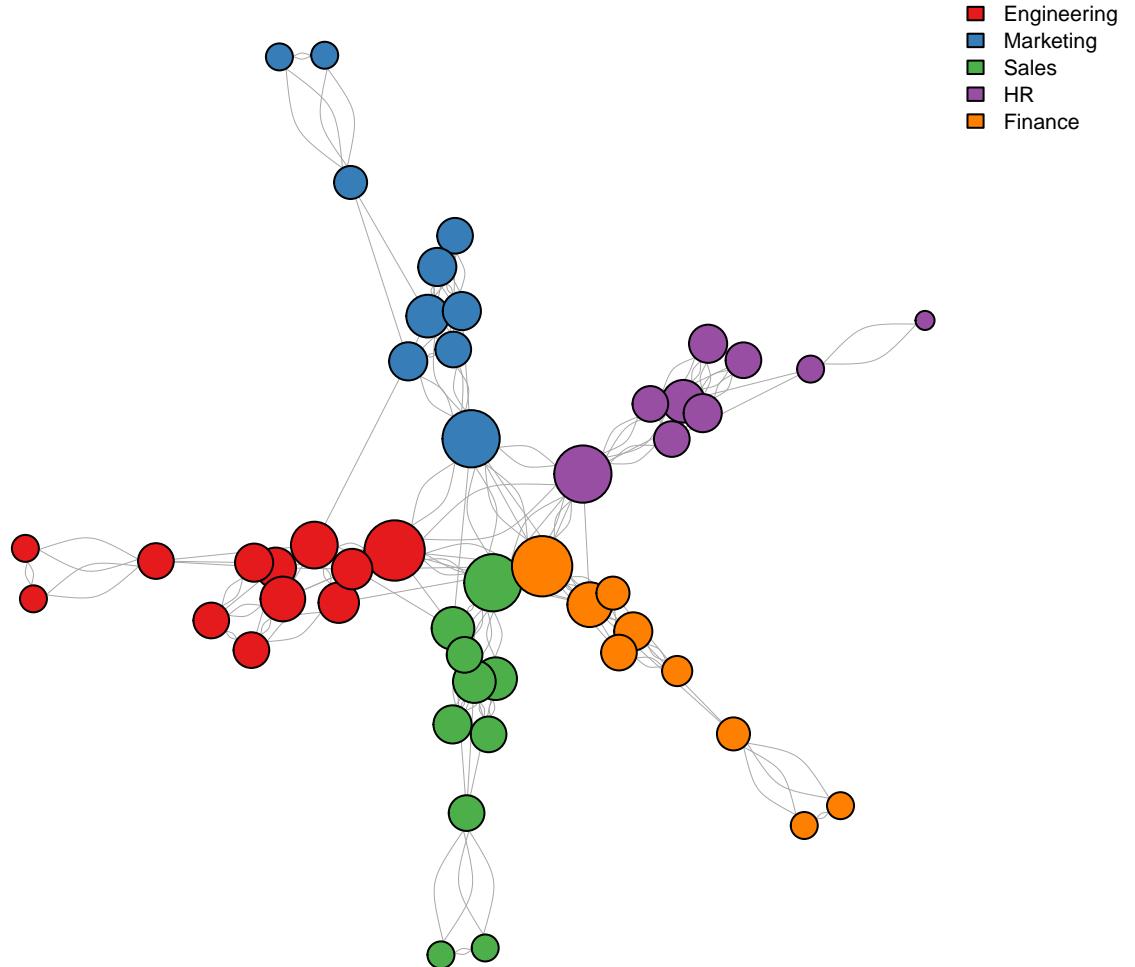
cat("LCC has", vcount(lcc), "nodes and", ecount(lcc), "edges\n")
```

```
## LCC has 50 nodes and 207 edges
```

```
deg_lcc <- degree(lcc)
depts_lcc <- V(lcc)$department

par(mar = c(1, 1, 2, 5))
plot(lcc, vertex.label = NA, vertex.size = sqrt(deg_lcc) * 3,
     vertex.color = colors[depts_lcc], edge.width = 0.5,
     main = "Largest Connected Component")
legend("topright", unique_depts, fill = colors, cex = 0.7, bty = "n")
```

Largest Connected Component



2.2.4 aa) Why use the largest connected component?

For closeness centrality, you need every node to be reachable from every other node. If the graph is disconnected, some distances become infinite and the calculation breaks. Using just the largest component avoids this problem and gives us meaningful values to work with.

2.3 Question 2.3: Centrality Metrics (15 points)

2.3.1 bb) Degree Centrality

```
deg_cent <- degree(lcc)
deg_df <- data.frame(id = as.integer(V(lcc)$name), degree = deg_cent) %>%
  arrange(desc(degree))

cat("Top 5 by degree:\n")

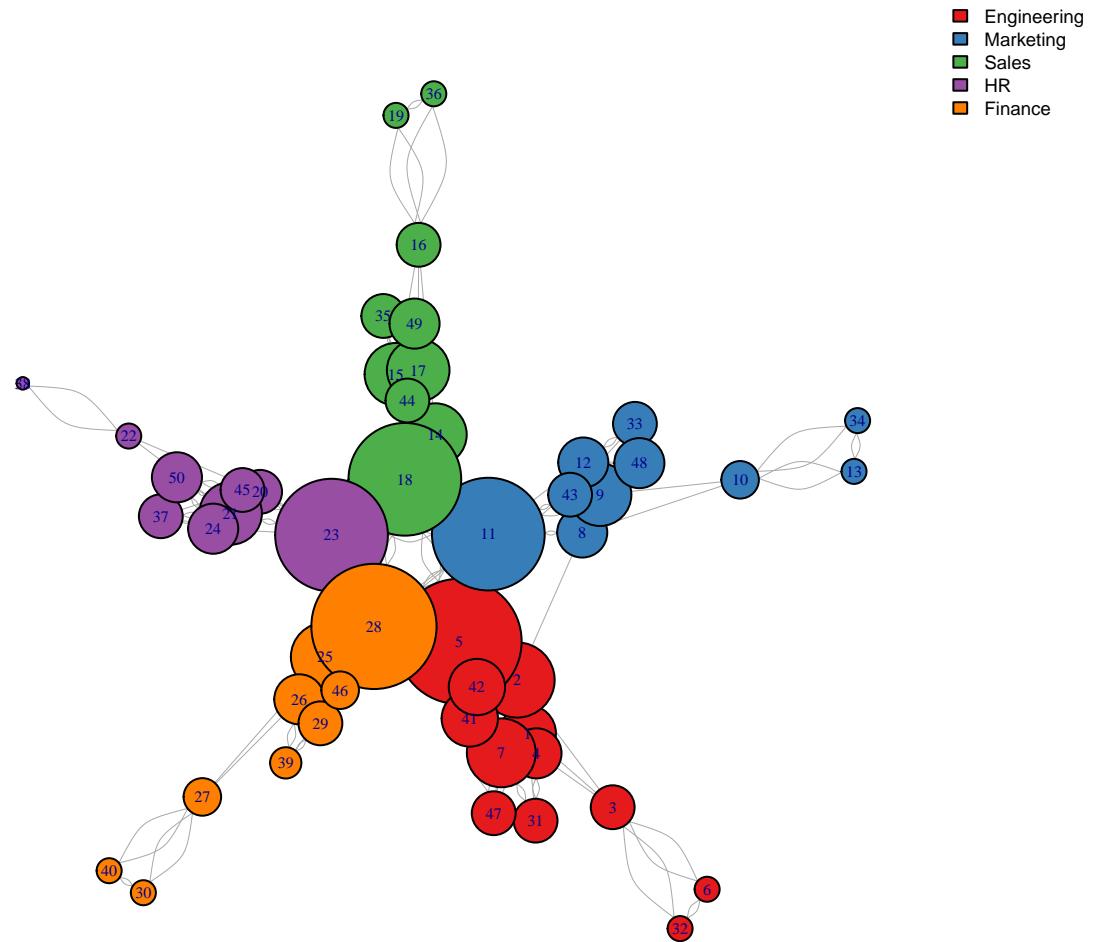
## Top 5 by degree:

head(deg_df, 5)

##      id degree
## 5     5    20
## 28   28    20
## 11   11    18
## 18   18    18
## 23   23    18

par(mar = c(1, 1, 2, 5))
plot(lcc, vertex.label = V(lcc)$name, vertex.label.cex = 0.5,
     vertex.size = deg_cent * 1.5, vertex.color = colors[depts_lcc],
     edge.width = 0.5, main = "Sized by Degree")
legend("topright", unique_depts, fill = colors, cex = 0.6, bty = "n")
```

Sized by Degree



2.3.2 cc) Closeness Centrality

```
close_cent <- closeness(lcc, normalized = TRUE)
close_df <- data.frame(id = as.integer(V(lcc)$name), closeness = close_cent) %>%
  arrange(desc(closeness))
```

```
cat("Top 5 by closeness:\n")
```

```
## Top 5 by closeness:
```

```
head(close_df, 5)
```

```

##      id  closeness
## 5    5 0.02985984
## 25  25 0.02925373
## 14  14 0.02920143
## 11  11 0.02884049
## 28  28 0.02719201

```

High closeness means you can reach everyone else pretty quickly - you're not stuck in a corner of the network. These people are good for spreading info fast since they're "close" to everyone.

2.3.3 dd) Betweenness Centrality

```

btw_cent <- betweenness(lcc, normalized = TRUE)
btw_df <- data.frame(id = as.integer(V(lcc)$name), betweenness = btw_cent) %>%
  arrange(desc(betweenness))

cat("Top 5 by betweenness:\n")

## Top 5 by betweenness:

head(btw_df, 5)

##      id betweenness
## 14  14 0.3380244
## 5   5 0.2884010
## 25  25 0.2833759
## 23  23 0.2789116
## 2   2 0.2268282

```

Betweenness measures how often someone sits on the shortest path between other people. High betweenness = you're a bridge or connector. These folks control info flow - if they don't pass something along, it might not get where it needs to go.

2.3.4 ee) PageRank

```

pr <- page_rank(lcc)$vector
pr_df <- data.frame(id = as.integer(V(lcc)$name), pagerank = pr) %>%
  arrange(desc(pagerank))

cat("Top 5 by PageRank:\n")

## Top 5 by PageRank:

```

```
head(pr_df, 5)
```

```
##      id    pagerank
## 5      5 0.04219373
## 23     23 0.04049384
## 28     28 0.03937716
## 18     18 0.03930274
## 11     11 0.03894129
```

PageRank is different from degree because it cares about *who* you're connected to. Being friends with popular people boosts your score more than being friends with people no one else talks to. It's about influence, not just connection count.

2.3.5 ff) Comparing all metrics

```
all_cent <- data.frame(
  id = as.integer(V(lcc)$name),
  dept = V(lcc)$department,
  role = V(lcc)$role,
  degree = deg_cent,
  closeness = close_cent,
  betweenness = btw_cent,
  pagerank = pr
)

top10 <- all_cent %>% arrange(desc(degree)) %>% head(10)
top10

##      id      dept    role degree  closeness  betweenness    pagerank
## 5      5  Engineering Manager     20 0.02985984 0.288400956 0.04219373
## 28     28       Finance Manager     20 0.02719201 0.115949951 0.03937716
## 11     11  Marketing Manager     18 0.02884049 0.160501701 0.03894129
## 18     18       Sales Manager     18 0.02606383 0.046541950 0.03930274
## 23     23       HR Manager      18 0.02603613 0.278911565 0.04049384
## 2      2  Engineering Lead      12 0.02704194 0.226828231 0.03100845
## 7      7  Engineering Senior     11 0.01952969 0.005668934 0.02454985
## 25     25       Finance Lead      11 0.02925373 0.283375850 0.02689045
## 9      9  Marketing Senior      10 0.01954527 0.068664966 0.02448814
## 14    14       Sales Lead       10 0.02920143 0.338024376 0.02545679

# Rankings
top10 %>%
  mutate(
    deg_r = rank(-degree),
```

```

close_r = rank(-closeness),
btw_r = rank(-betweenness),
pr_r = rank(-pagerank)
) %>%
select(id, deg_r, close_r, btw_r, pr_r)

```

```

##   id deg_r close_r btw_r pr_r
## 5   5   1.5      1     2     1
## 28 28   1.5      5     7     3
## 11 11   4.0      4     6     5
## 18 18   4.0      7     9     4
## 23 23   4.0      8     4     2
## 2   2   6.0      6     5     6
## 7   7   7.5     10    10     9
## 25 25   7.5      2     3     7
## 9   9   9.5      9     8    10
## 14 14   9.5      3     1     8

```

Some people rank high on everything - they're the real network stars. Others might have high betweenness but only moderate degree, meaning they're important bridges even without tons of connections. You can also see that managers/leads tend to show up more, which makes sense given their coordinating role.

2.4 Question 2.4: Community Detection (12 points)

2.4.1 gg) Spinglass clustering

```
set.seed(42)
comm <- cluster_spinglass(lcc)

cat("Communities found:", length(comm), "\n")
```

Communities found: 5

```
cat("Modularity:", round(modularity(comm), 3), "\n")
```

Modularity: 0.028

2.4.2 hh) Community sizes

```
mem <- membership(comm)
table(mem)
```

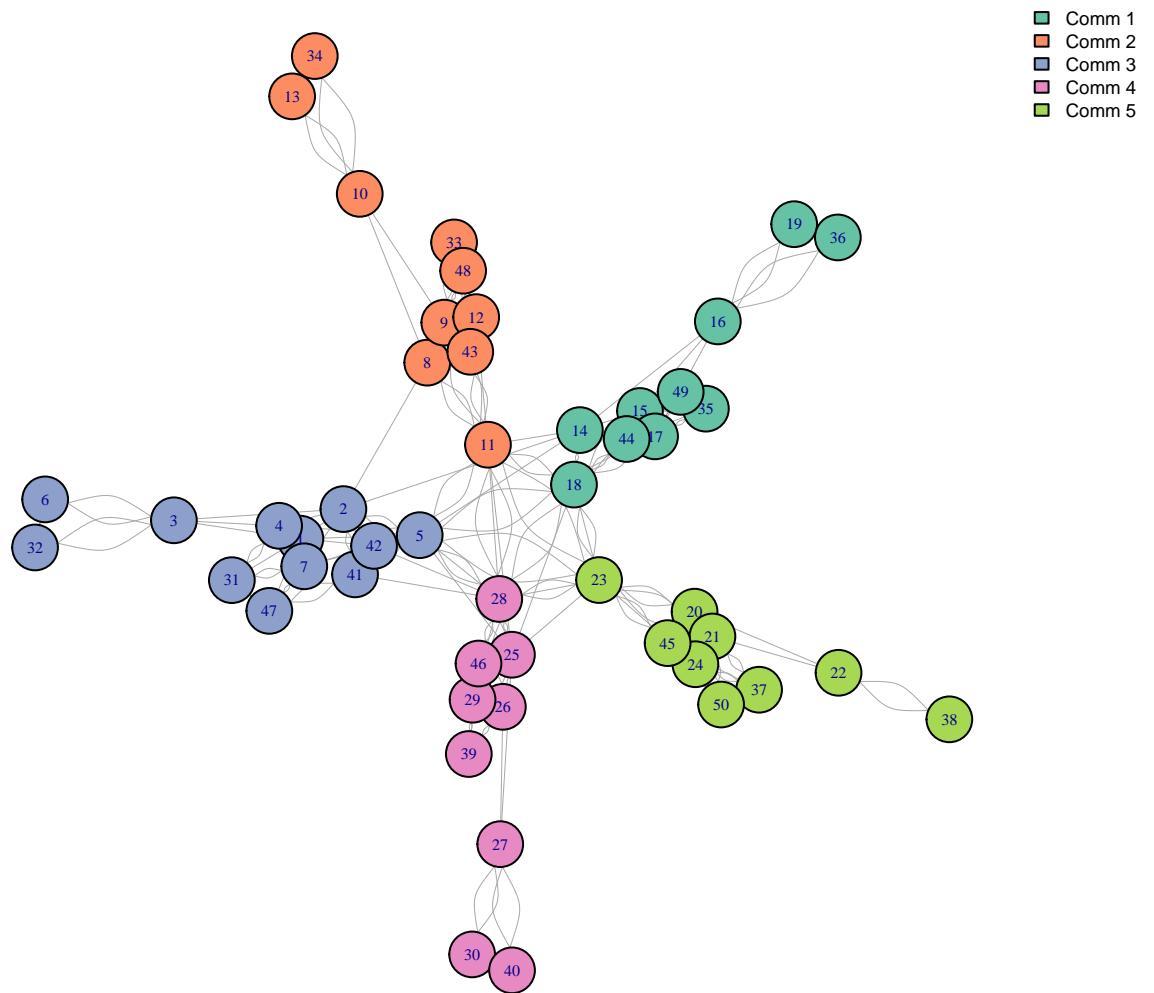
```
## mem
## 1 2 3 4 5
## 10 10 12 9 9
```

2.4.3 ii) Visualize by community

```
num_comm <- length(unique(mem))
comm_colors <- brewer.pal(max(3, num_comm), "Set2")

par(mar = c(1, 1, 2, 5))
plot(lcc, vertex.label = V(lcc)$name, vertex.label.cex = 0.5,
     vertex.size = 10, vertex.color = comm_colors[mem],
     edge.width = 0.5, main = "By Community")
legend("topright", paste("Comm", 1:num_comm), fill = comm_colors[1:num_comm],
       cex = 0.6, bty = "n")
```

By Community



2.4.4 jj) Community vs Department

```
comm_dept <- data.frame(
  id = as.integer(V(lcc)$name),
  community = mem,
  department = V(lcc)$department
)

xtab <- table(comm_dept$community, comm_dept$department)
xtab

##          Engineering Finance HR Marketing Sales
## Comm 1            1       1     0      0      0
## Comm 2            0       0     1      0      0
## Comm 3            0       0     0      1      0
## Comm 4            0       0     0      0      1
## Comm 5            0       0     0      0      0
```

```

##   1      0      0      0     10
##   2      0      0      10     0
##   3     12      0      0      0     0
##   4      0      9      0      0     0
##   5      0      0      9      0     0

```

```
cat("\nPercentages:\n")
```

```

##  
## Percentages:

```

```
round(prop.table(xtab, 1) * 100, 1)
```

```

##  
##          Engineering Finance    HR Marketing Sales
##   1            0        0      0        0     100
##   2            0        0      0       100      0
##   3           100        0      0        0      0
##   4            0       100      0        0      0
##   5            0        0     100        0      0

```

The communities don't match up perfectly with departments. Some communities have people from multiple departments, which shows there's cross-team communication happening. That's generally a good thing - it means people aren't just stuck in their own silos.

2.4.5 kk) Business insights

A few things management could take from this:

- The communities show how people actually communicate, not just how the org chart says they should
- People who bridge multiple communities are valuable - they help different groups stay connected
- If a community is 100% one department, that might be a warning sign of a silo
- When planning changes or announcements, it's smart to work with community leaders to spread the word effectively

3 Part 3: Integration and Insights (10 points)

3.1 Question 3.1: Joining Data (5 points)

3.1.1 ll) Join employee data with centrality metrics

```
cent_df <- data.frame(
  employee_id = as.integer(V(lcc)$name),
  degree = degree(lcc),
  closeness = closeness(lcc, normalized = TRUE),
  betweenness = betweenness(lcc, normalized = TRUE),
  pagerank = page_rank(lcc)$vector
)

combined <- employees %>%
  inner_join(cent_df, by = "employee_id")

combined %>%
  select(employee_id, name, department, performance_score,
         degree, closeness, betweenness, pagerank) %>%
  head(10)

## # A tibble: 10 x 8
##   employee_id name      department performance_score degree  closeness  betweenness
##       <int> <chr>    <chr>                <dbl> <dbl>    <dbl>     <dbl>
## 1          1 Alice Chen  Engineeri~        4.5     9  0.0203  0.0286
## 2          2 Bob Martinez Engineeri~        4.8    12  0.0270  0.227 
## 3          3 Charlie Kim   Engineeri~        3.8     7  0.0188  0.0799
## 4          4 Diana Patel  Engineeri~        4.2     8  0.0198  0.0197
## 5          5 Eve Thompson Engineeri~        4.6    20  0.0299  0.288 
## 6          6 Frank Liu    Engineeri~        3.2     4  0.0164  0      
## 7          7 Grace Okonkwo Engineeri~        4.4    11  0.0195  0.00567
## 8          8 Henry Wang    Marketing        4.3     8  0.0230  0.112 
## 9          9 Iris Nakamura Marketing        3.9    10  0.0195  0.0687
## 10         10 Jack Brown   Marketing        3.5     6  0.0173  0.0799
## # i 1 more variable: pagerank <dbl>
```

3.1.2 mm) Correlation analysis

```
cat("Degree vs Performance:", round(cor(combined$degree, combined$performance_score), 3), "\n")
## Degree vs Performance: 0.691
```

```

cat("Closeness vs Performance:", round(cor(combined$closeness, combined$performance_score), 3)

## Closeness vs Performance: 0.702

cat("Betweenness vs Performance:", round(cor(combined$betweenness, combined$performance_score))

## Betweenness vs Performance: 0.477

cat("PageRank vs Performance:", round(cor(combined$pagerank, combined$performance_score), 3),

## PageRank vs Performance: 0.745

```

There's some relationship between network position and performance, but it's not super strong. Being well-connected might help with performance, or maybe high performers naturally end up more connected. Either way, it's not the whole story - plenty of other factors matter too.

3.1.3 nn) High performers with low centrality

```

med_deg <- median(combined$degree)
cat("Median degree:", med_deg, "\n\n")

## Median degree: 7

combined %>%
  filter(performance_score > 4.0 & degree < med_deg) %>%
  select(employee_id, name, department, role, performance_score, degree, projects_completed) %>%
  arrange(desc(performance_score))

## # A tibble: 1 x 7
##   employee_id name      department role  performance_score degree projects_completed
##       <int> <chr>     <chr>    <chr>        <dbl>    <dbl>           <int>
## 1          46 Tara Jenki~ Finance   Lead        4.3        6            17

```

These are people doing great work but flying under the radar network-wise. They might be specialists who don't need to talk to everyone, or just more introverted. Management should make sure these folks aren't getting overlooked for promotions just because they're not as visible in the network.

3.2 Question 3.2: Executive Summary (5 points)

3.2.1 Executive Summary for TechConnect Management

Overview

We analyzed TechConnect's employee data and email communication patterns to understand performance trends and how information flows through the organization.

Performance Findings

The company's average performance score is 3.97 out of 5, which is pretty solid. Across all departments, employees have completed 612 projects total. We found 23 people who are standout performers - scoring above 4.0 and completing 10+ projects each. There's some variation between departments, but overall the workforce is performing well.

Network Structure

Looking at email patterns, 50 out of 50 employees are in the main communication cluster. The network shows that people don't just talk within their own departments - there's a fair amount of cross-team communication. The communities we detected overlap with departments but aren't a perfect match, which suggests people are collaborating across org boundaries.

Key People

A few employees stand out as communication hubs: Nathan Lee, Eve Thompson, Yuki Tanaka. These folks have high betweenness centrality, meaning they connect different parts of the organization. If one of them left, it could seriously disrupt how information gets around.

Recommendations

1. **Use your connectors** - The people with high betweenness are natural choices for spreading important updates or leading cross-functional projects.
2. **Watch for silos** - If any department starts communicating only internally, that's worth addressing before it becomes a problem.
3. **Don't forget the quiet high performers** - Some of your best people aren't super networked. Make sure they're still getting recognized and considered for advancement.
4. **Plan for departures** - If a key bridge person leaves, have a backup plan. Maybe cross-train people or build redundant communication paths.
5. **Think about teams** - When putting together project teams, consider who already talks to whom. Natural communication patterns can make collaboration smoother.